COMP 50001: Coursework 1

Autumn Term, 2024

1 Introduction

The goal of this coursework is to implement algorithms that will be used to control your army in a game of IMPERIAL CONQUEST: a real-time strategy game created for this course. It takes place in a galaxy far, far away, where space travel to different planets is made possible through a network of wormholes. Planets are able to produce fleets of ships that can be sent to divide and conquer the enemy.


Figure 1: IMPERIAL CONQUEST boasts a fashionable text-based interface (though it will only come to life for the second part of the coursework)

Both coursework assignments will use the same infrastructure. This first assignment concerns itself with choosing a basic initial strategy and some simple pathfinding. The second assignment will deal with more complex strategies and dealing with opponents.

2 Structure

Like your Haskell PPTs from last year, this coursework follows a standard Haskell project structure using cabal:

```
comp50001_cw1_username
└── src
    ├── AdjList.hs ...................................... please complete
    ├── Examples.hs ..... Examples you might find useful for testing
    └── Game.hs .............................. Core datatypes, do not edit
```

The various problems outlined in this specification will be completed in the `Optimise.hs`, `ShortestPaths.hs`, `Heap.hs`, and `AdjList.hs` files. The automated testing will replace the contents of all files marked *do not edit*, as well as the tests and project `imperial-conquest.cabal` files.

During development, you might find it useful to load the source file into `GHCi` by running `cabal repl`. This way, you can test your solutions, and ensure there are no compilation errors as you're working on your code.

This coursework will be automarked: test suites are provided to check your answers. To run the tests, you should run the following command in the root directory of the repository:

```
cabal test
```

The submission of your coursework should be done via LabTS using `git` commits in the usual way.

3   Background

IMPERIAL CONQUEST is a game inspired by Galcon, a two player real-time strategy game. A game of IMPERIAL CONQUEST is played between two players, and time progresses in a series of turns.

The `Player` data type represents the two players of the game.

```
data Player = Player1 | Player2
```

This first coursework involves no interaction between players, but the data types will be shared between the second coursework, and so are defined here.

The game is played on a map with a number of planets. Planets are represented by the `Planet` type.

```
data Planet = Planet Owner Ships Growth
```

All of the datatypes in this section are found in `Game.hs` and are relied upon for communication with the server and should not be modified.

A `newtype` uses an existing type as the basis of a new type by wrapping values in a new constructor. Here it means that a `Ship` cannot be confused with a `Growth`, even though both are fundamentally storing an `Int`.

```haskell
newtype Ships = Ships Int
newtype Growth = Growth Int
```

Thus, a planet is a value `Planet` owner ships growth, where an owner can be either neutral, or one of the two players:

```haskell
data Owner = Neutral | Owned Player
```

Planets have a number of spaceships in garrison, represented by ships that belong to the owner of the planet. Finally, each turn, planets that are owned by a player have factories which can produce a number of ships given by the growth value.

In order to refer to planets, they each have an identifying number, which is presented as a value of type `PlanetId`:

```haskell
newtype PlanetId = PlanetId Int
```

The `Planets` type represents a collection of planets as a key-value map from planet identifiers to the planet structures:

```haskell
type Planets = Map PlanetId Planet
```

A `Map` is a data structure that is provided in the `Data.Map` module, which is documented at `https://hackage.haskell.org/package/containers-0.7/docs/Data-Map-Lazy.html`, which gives operations and their complexities.

The planets in this galaxy are too far away for spaceships to travel between them within a reasonable amount of time using traditional propulsion systems, but certain planets are connected by wormholes, allowing spaceships to travel through them faster than light.

```haskell
data Wormhole = Wormhole Source Target Turns

newtype Source = Source PlanetId
newtype Target = Target PlanetId
newtype Turns  = Turns Int
```

Crucially, a `Wormhole` connects two planets in a *directed* way. That is, wormholes are one-way streets with a value of type `Source` for the source, and `Target` for the target. The value of type `Turns` indicates how many turns it takes to travel through the wormhole.

Similarly to `Planets`, `Wormhole`s are also referred to by an identifier which is just a wrapper around an `Int`:

```haskell
newtype WormholeId = WormholeId Int
```

These are collected into a key-value map in much the same way:

```haskell
type Wormholes = Map WormholeId Wormhole
```

As spaceships go through wormholes, the number of turns left on their journey before they arrive is tracked. The `Fleet` data type represents a fleet of ships by storing whose ships they are, how many of them are there, which wormhole they're in, and how many turns left before they reach their target.

```
data Fleet = Fleet Player Ships WormholeId Turns
```

When a fleet of ships arrives at the target of the wormhole, they join forces with the ships that are there if they belong to the same owner. Otherwise, the fleet attacks the ships on the defending planet, cancelling each other out one-to-one. If there are more ships in the fleet than on the planet then the remaining ships from the fleet capture the planet garrison, and the planet is owned by the fleet owner.

The `Fleets` type represents a collection of fleets. Note that these do not have identifiers, as they are not referred to anywhere.

```
type Fleets = [Fleet]
```

The state of the whole game at any given time is represented by the `GameState` data type:

```
data GameState = GameState Planets Wormholes Fleets
```

Finally, at every turn, the server takes a list of orders and begins to execute them.

```
data Order = Order WormholeId Ships
```

The exact rules of the game are not important for this assignment, though they will become relevant in the next.

4 Dynamic Programming

Dynamic programming is a technique for optimising the runtime performance of a recursive algorithm that has overlapping subproblems. The speedup comes from storing the subsolutions for later use instead of recomputing them every time they are needed.

The strategy is developed in two stages:

1. Write an inefficient recursive algorithm that solves the problem.

2. Improve efficiency by storing intermediate shared results.

As a simple example, dynamic programming can be applied to speed up a naive implementation of the Fibonacci function.

Each Fibonacci number for a given value n is given by `fib n` given by the following recursive algorithm:

While the input to `fib` is usually bound by an `Int`, even for small inputs the function will overflow in its output, hence the return type is `Integer`, which can represent arbitrarily large integers.

```
fib :: Int -> Integer
fib 0 = 0
fib 1 = 1
fib n = fib (n-2) + fib (n-1)
```

Figure 2: Call tree of fib 10, showing multiple repeated computations.

This code is remarkably inefficient, since there are repeated calls to computations that recalculate the same value. Figure 2 shows that `fib 8` is called twice in the calculation of `fib 10`. In turn, `fib 7` is called three times, `fib 6` five times, etc. As a lower bound of its complexity, we have the following recurrence relation:

$$T_{\texttt{fib}}(n) \geq 1 + 2 \times T_{\texttt{fib}}(n-2)$$

Solving this recurrence gives $T_{\texttt{fib}}(n) \in O(2^{n/2})$. This is remarkably expensive and due to the fact that there are large overlaps in the solutions of subproblems, which keep getting recalculated.

This recalculation can be avoided by building up the intermediate values up to the result. More concretely, the computation of `fib n`, involves defining an array containing the values of `fib` for *all* numbers from 0 up to n.

```haskell
table :: Int -> Array Int Integer
table n = array (0,n) [ (0, 0)
                      , (1, 1)
                      , (2, table ! 0 + table ! 1)
                      , (3, table ! 1 + table ! 2)
                      , ... ]
```

The `Array Int Integer` type represents an array indexed by `Int` containing values of type `Integer`.

The function (`!`) provides constant-time random access.

The function `array` takes a range `(u,v)` of values as its bounds, and a list of pairs where each pair `(i,x)` is used to place `x` at index `i` in the table.

Notice that every element of `table` refers to solutions of previous problems that eventually lead to a base case: there are no circular references so the self-referential definition is well-defined. Here is a version of `fib` that builds the right `table` and returns the last element of the array.

This is, of course, not the best way to calculate Fibonacci numbers (which can be done in sublinear time), but it illustrates how a recursive algorithm can be made more efficient.

```haskell
fib' :: Int -> Integer
fib' n = table ! n
  where
    table :: Array Int Integer
    table = tabulate (0, n) mfib

    mfib 0 = 0
    mfib 1 = 1
    mfib n = table ! (n-1) + table ! (n-2)
```

It is important that `mfib` is in the same level of scope as `table`: if it were top-level then a new table would be created on each call!

The table given by `tabulate (x,y) f` contains the results of applying `f` to all the values between `x` and `y`. It is implemented as an array which gives constant time access to its elements. The function `tabulate` can be found in `Util.hs`:

```haskell
tabulate :: Ix i => (i,i) -> (i -> a) -> Array i a
tabulate (u,v) f =
  array (u,v) [ (i, f i) | i <- range (u, v)]
```

The constraint `Ix i` allows the values of type `i` to be drawn from those given by the `range` function, as well as enabling values to be indexed over in an array. This allows arrays to be indexed by types other than `Int`. For instance, a valid index is a tuple `(Int, Int)` for a two-dimensional array indexed by pairs of `Int`s.

The cost of building this table is the sum of all the individual calls of f. The key to efficiency is that the function f can itself refer to the table that is being constructed. If the cost of f is constant, such that $T_f(i) \in O(1)$, and the table has n elements, then the cost of its construction is $T_{table}(n) \in O(n)$.

In the code above, mfib is a local version of fib that finds values in the table rather than by recursion. The function takes constant time since (!) is a constant time operation. Thus, the time complexity of evaluating fib' n is given by:

$$T_{fib'}(n) = 1 + T_{table}(n) + T_!(n)$$

where $T_{table}(n)$ is the time it takes to construct the table, and $T_!(n)$ is the time it take to look up a value in that table. Therefore, the overall cost is $T_{fib'}(n) \in O(n)$, which is much better than before.

The slogan for dynamic programming algorithms is to *memoize shared sub-computations*: by memoizing the results of sub-computations in a table, repeated computations are avoided whenever possible, resulting in a much faster algorithm.

5   Planet Picking

The first task of this coursework is to plan which planets to conquer first. To simplify things, assume that there are a number of planets that are equally reachable from your forces.

As an example, here is a map with five planets, planet 0 being your home planet, and planets 1..4 being the neutral planets that you can conquer, each in a single turn. Furthermore, the planets are not reachable from one another.

```
example1 :: GameState
example1 = GameState planets wormholes fleets where
  planets = M.fromList
    [ (PlanetId 0, Planet (Owned Player1) (Ships 300) (Growth 0))
    , (PlanetId 1, Planet Neutral        (Ships 200) (Growth 50))
    , (PlanetId 2, Planet Neutral        (Ships 150) (Growth 10))
    , (PlanetId 3, Planet Neutral        (Ships 30)  (Growth 5))
    , (PlanetId 4, Planet Neutral        (Ships 100) (Growth 20))
    ]
  wormholes = M.fromList
    [ (WormholeId 0, Wormhole homePlanet (Target 1) (Turns 1))
    , (WormholeId 1, Wormhole homePlanet (Target 2) (Turns 1))
    , (WormholeId 2, Wormhole homePlanet (Target 3) (Turns 1))
    , (WormholeId 3, Wormhole homePlanet (Target 4) (Turns 1))
    ] where homePlanet = Source 0
  fleets = []
```

The information in a `GameState` can be queried. For instance, `Game.hs` provides both the `targetPlanets` function, which lists the planets that can be reached from a given source, and the `shipsOnPlanet` function, which finds how many ships are garrisoned on a planet:

```haskell
targetPlanets :: GameState -> Source -> [(PlanetId, Ships, Growth)]
targetPlanets st s =
  map (planetDetails . target) (M.elems (wormholesFrom s st))
  where
    planetDetails :: PlanetId -> (PlanetId, Ships, Growth)
    planetDetails pId = (pId, ships, growth)
      where Planet _ ships growth = lookupPlanet pId st


shipsOnPlanet :: GameState -> PlanetId -> Ships
shipsOnPlanet st pId = ships
  where Planet _ ships _ = lookupPlanet pId st
```

Both of these functions use `lookupPlanet` to extract a planet from the `GameState`:

```haskell
lookupPlanet :: PlanetId -> GameState -> Planet
lookupPlanet pId (GameState ps _ _) = fromJust (M.lookup pId ps)
```

It is also possible to determine the wormholes that correspond to a planet, whether that be wormholes to or from that planet:

```haskell
wormholesFrom :: Source -> GameState -> Wormholes
wormholesFrom pId (GameState _ ws _) =
  M.filter (\(Wormhole s _ _) -> s == pId) ws


wormholesTo :: Target -> GameState -> Wormholes
wormholesTo pId (GameState _ ws _) =
  M.filter (\(Wormhole _ t _) -> t == pId) ws
```

Here is an example of how `cabal repl` can be used to query some values:

```haskell
ghci> targetPlanets example1 (Source (PlanetId 0))
[(1,200,50),(2,150,10),(3,30,5),(4,100,20)]
```

If the overall number of defenders among all these planets is overwhelming, you will have to pick which planets to conquer first. Planets have a different intrinsic value, which is represented by their growth rate. The question is *which planets to conquer in the first turn to maximise the growth rate obtained in this turn*. In the case of `example1`, the optimal strategy if your fleet has 300 ships is to conquer planet 1 and planet 4, resulting in a growth rate of 70.

Planets have different intrinsic value, so you should aim to maximise the value of your conquest given your capacity to attack. This is optimisation problem is an example of the classic *knapsack* problem.

5.1   Unbounded Knapsack

The unbounded knapsack problem concerns itself with packing a knapsack with some given capacity `c` with elements of some weight and value drawn from a list. There is no limit to the number of times an element can be picked from the list. The goal is to maximise the value of the items that can be placed in the list, without going over the capacity.

Thus, the available items are represented as a list of triples, where each item has a name, some weight, and some value.

```
[(name, weight, value)]
```

The name parameter is not used in this first function, but it will become useful later.

As a recursive algorithm, `knapsack` is easily stated: take an item (`name`, `weight`, `value`), and consider what would happen if we put it in the knapsack. The total value would be increased by the `value` of the item, but the remaining capacity would be decreased by the `weight` of the item. So, if the starting capacity is `c`, then picking this item would result in a maximum value of the item's `value` plus the maximum value of the subproblem where the capacity is `c-w` (since after putting this item in, we need to maximise the value by filling in the remaining capacity). The optimal solution is achieved by trying every item as the first one, then recursively optimising the remaining capacity, ultimately finding the maximum one.

The `knapsack` function is highly polymorphic, as it works on any types of inputs, as long as they can be compared (as stated by the `Ord` constraint) and support basic arithmetic (as stated by the `Num` constraint). Here the abstraction helps make clear exactly which properties of the types are relied upon.

This recursive description can be written as the following:

```
knapsack :: forall name weight value .
        (Ord weight, Num weight, Ord value, Num value) =>
        [(name, weight, value)] -> weight -> value
knapsack wvs c = maximumElse 0 [ v + knapsack wvs (c - w)
                        | (_,w,v) <- wvs , w <= c ]
```

That is, from the input list `wvs`, find all items whose weight is less than the capacity, and try them by recursively solving the smaller problem. Finally, take the largest element of this list using `maximumElse`, found in `Util.hs`:

```
maximumElse :: Ord a => a -> [a] -> a
maximumElse x xs = foldr max x xs
```

The maximum value here is bounded below by `0`.

This implementation of `knapsack` correctly produces the maximum value, which can be tested with `cabal repl`:

```
ghci> knapsack [("a", 35, 10), ("b", 153, 200), ("c", 100, 20)] 800
1010
```

However, the efficiency of this algorithm is quite terrible when the knapsack can be filled to a given capacity by picking different combinations of elements, since this leads to repeated computations. Try increasing the capacity to see how slow it gets.

The `forall` is used to put the type variables `name`, `weight`, and `value` in scope so that they can be referred to in the `where` clause.

Notice the additional (`Ix` weight) type class constraint in the function's signature. This indicates that the `weight` will be used as an array index when building

Problem 1: *Dynamic Knapsack*

Use dynamic programming to improve the running time of `knapsack`, by giving a definition of `mknapsack` in `Optimise.hs`.

```haskell
knapsack' :: forall name weight value .
  (Ix weight, Num weight, Ord value, Num value) =>
  [(name, weight, value)] -> weight -> value
knapsack' wvs c = table ! c
  where
    table :: Array weight value
    table = tabulate (0,c) mknapsack

    mknapsack :: weight -> value
    mknapsack c = undefined
```

HINT: compare the recursive version of `fib` with the dynamic programming version to see how to translate one to the other.

While this correctly and efficiently calculates the maximum value of the knapsack, it does not announce what the items that are picked actually are (in other words, the `name` elements are ignored). To do this, the return type `Value` needs to be modified to something that holds both the value and the index of the item that was chosen. Then, you will have to modify the algorithm so that the index does not get in the way, and so that the indices are properly combined.

Problem 2: *Knapsack Elements*

Implement `knapsack''` as a modified version of `knapsack'` so that it outputs the maximum value of the knapsack, as well as a list of the element indices that are chosen to obtain that value.

```haskell
knapsack''
  :: forall name weight value .
  (Ix weight, Num weight, Ord value, Num value) =>
  [(name, weight, value)] -> weight -> (value, [name])
knapsack'' wvs c = table ! c
  where
    table :: Array weight (value, [name])
    table = tabulate (0,c) mknapsack

    mknapsack :: weight -> (value, [name])
    mknapsack c = undefined
```

> HINT: You may need to make use of `maximumBy` to be able to ig-
> nore the indices that are being used. A fantastic site for finding
> useful functions is `https://hoogle.haskell.org/`, where even type
> signatures can be typed in.

If your solution is correct, you should see something like this:

```
ghci> knapsack'' [("a", 35, 10), ("b", 153, 200), ("c", 100, 20)] 800
(1010,["b","b","b","b","b","a"])
```

5.2 *Bounded Knapsack*

The unbounded knapsack problem allows the items to be used multiple
times. However, the planets can only be conquered once, so it doesn't
make much sense to allow them to be picked multiple times. The goal
of this section is to implement the *bounded knapsack problem* to help
decide which planets should be conquered.

Problem 3: *Bounded Knapsack*

This is the type of `bknapsack`, which is similar to `knapsack` except
that the elements are not replaced when picked.

```
bknapsack :: forall name weight value .
  (Ord weight, Num weight, Ord value, Num value) =>
  [(name, weight, value)] -> weight -> (value, [name])
```

HINT: Perform case analysis on the list of items. After using an item,
remove it from the recursive call.

To implement this function, first write a *recursive* definition, where
the result of `bknapsack wvs c` is the maximum value that can be
packed into a capacity of `c` from the elements of `wvs` by only using
each element at most once.

The solution to `bknapsack` (hopefully) looks simpler, or at least
not much more difficult, than for `knapsack`. However, it turns out
to be trickier to apply dynamic programming to the bounded case
than the unbounded case. To understand why, consider the previous
applications of dynamic programming. Both in the case of `fib` and
`knapsack`, the solution was to inductively iterate through the space of
inputs until reaching the desired value. In case of `fib`, the `Int`s were
iterated, and in the case of `knapsack`, the `weight` values. The reason
this approach works for `knapsack` is because in the recursive call, only
the `weight` parameter changes, so the table can be built only indexed
by that parameter alone.

In the case of bknapsack, there are two varying parameters: the weight as before, but also the input list, as the picked element gets removed in every iteration.

A first approach might be to simply make the array indexed by not only the weight parameter, but also the [(name, weight, value)] list. However, indexing on this list of triples would be a very large domain, which would make it difficult to create an efficient lookup table.

Problem 4: *Bounded Knapsack Revisited*

Write another recursive solution, bknapsack', to the bounded knapsack problem. This time, do not change the list in the recursion, but instead introduce another parameter that keeps track of progress. That is, bknapsack' wvs i will use only the first i (counting from 1) items in the list wvs.

```
bknapsack' :: forall name weight value .
  (Ord weight, Num weight, Ord value, Num value) =>
  [(name, weight, value)] -> Int ->
  weight -> (value, [name])
```

HINT: Take a careful look at the type signature, and notice the extra Int parameter. Use this to keep track of which part of the list has been processed.

Problem 5: *Dynamic Bounded Knapsack*

Using dynamic programming implement bknapsack'', which is the efficient version of bknapsack' that makes use of tables rather than repeated recursion.

```
bknapsack'' :: forall name weight value .
  (Ord name, Ix weight, Ord weight, Num weight,
    Ord value, Num value) =>
  [(name, weight, value)] -> weight -> (value, [name])
```

HINT: when populating the table, use the new Int parameter instead of the input list as part of the index.

HINT: The Array type can be indexed by any type that has an Ix instance. Notably, tuples of indexable types are also indexable, which means for example it is possible to have an array of type Array (Int, Int) Bool, which can be thought of as a two dimensional array.

To put all the pieces together, the optimal conquering strategy can

finally be calculated from any source planet:

```
optimise :: GameState -> Source -> (Growth, [PlanetId])
optimise st s@(Source p) =
    bknapsack'' (targetPlanets st s) (shipsOnPlanet st p)
```

For example, the example1 map state should yield

```
ghci> optimise example1 (Source 0)
(70,[4,1])
```

6   Directed Graphs

Navigating through the network of wormholes will require a representation of that network, and this is nicely achieved by using a *graph*. A graph consists of vertices and edges. For simplicity, we will focus on the definition of a weighted directed simple graph: since it is weighted and directed, each edge has a specified source, target, and weight. Weights will be assumed to simply be of type Integer. All the types in this section are located in Graph.hs.

There are many kinds of graphs, depending on whether there are cycles, whether edges are directed, whether there can be only up to one edge between to vertices, and whether the edges have weights.

```
type Weight = Integer
```

This is encapsulated by the following class, where e is the type of edges, and v is the type of vertices:

The e -> v part of the class declaration is called a *functional dependency*, and it describes a relation between edges and vertices. Here it says that knowing the type of an edge determines the type of the vertex.

```
class Eq v => Edge e v | e -> v where
  source :: e -> v
  target :: e -> v
  weight :: e -> Weight
```

This abstraction means that different types can be considered to be edges, which will become useful later on.

As an example, a triple (String, String, Integer) can be thought of as an edge between vertices that are Strings.

```
instance Edge (String, String, Integer) String where
  source (s, _, _) = s
  target (_, t, _) = t
  weight (_, _, i) = i
```

Concretely, ("here","there",10) is an edge from the source "here" to the target "there" with weight 10.

For IMPERIAL CONQUEST, the Wormhole type can readily be made an instance of Edge with vertices of type PlanetId.

```
instance Edge Wormhole PlanetId where
  source (Wormhole (Source s) _ _)    = s
  target (Wormhole _ (Target t) _)    = t
  weight (Wormhole _ _ (Turns turns)) = toInteger turns
```

Here, the weight of the edge is the number of turns it takes to travel through the wormhole.

Describing edges in this abstract sense using a type class has some nice benefits. It makes it possible to talk about algorithms on graphs in a way that is not tied to any specific representation. The advantages of abstraction here are twofold: the implemented algorithms are easier to understand because they do not refer to low-level details, and they are also more reusable. Just like knapsack, the graph algorithms here will also work on our galaxy, even though the implementations do not mention anything about planets or ships.

It is also possible to attach additional payload to edges when convenient. For instance, since a `Wormhole` is an edge, so is a pair of a `WormholeId` and a `Wormhole`.

This is convenient because when talking about the wormholes in the galaxy, it will be useful to know the identifier as well as the value.

```
instance Edge (WormholeId, Wormhole) PlanetId where
  source (_, w) = source w
  target (_, w) = target w
  weight (_, w) = weight w
```

An edge connects two points in a graph, and is a primitive building block of a graph. Edges compose together when the target of one agrees with the source of the other, making it possible to follow a path from a source vertex to a target vertex through a series of connected edges. In this sense, a *path* between two vertices can be thought of as a list of edges. Here is a concrete implementation of a `Path`:

```
data Path e = Path Weight [e]
```

As an optimisation, it is convenient for a path to also keep track of the sum of all the weights of its edges as an additional parameter. This is redundant information since the total weight could always be recomputed from just the list alone, but it makes it easier to inspect such a useful value.

An edge can trivially be turned into a path:

```
pathFromEdge :: Edge e v => e -> Path e
pathFromEdge e = Path (weight e) [e]
```

This is a path with just one edge.

Since the `Path` type contains a list, it is a lot more efficient to prepend a new element at the front (since (`:`) is $O(1)$) than it is to append at the end (which will cost $O(n)$ to traverse). A path in general can grow on either end (an edge going out of the target, or an edge going into the source), so there is a choice of which operation should be supported more efficiently.

There are of course other data structures such as double-ended queues that support both operations in amortized constant time.

For the purposes of this coursework, a fixed source for paths will be more useful, so a path will only be extended by adding edges out of the

target. To support this operation efficiently, the source will be stored at the tail end of the list, and the target is at the head. This means that a path can be extended by with an edge going out of the current target of the path.

```haskell
extend :: Edge e v => Path e -> e -> Path e
extend (Path _ []) _ = error "extend: Empty path"
extend (Path d (e:es)) e'
  | target e == source e' = Path (d + weight e') (e':e:es)
  | otherwise = error "extend: Incompatible endpoints"
```

Given a list of edges es = [e_0 , ..., e_n] where each e_i has source vertex v_i, target vertex v_{i+1}, and weight w_i, they can be stitched together into a path:

```haskell
pathFromEdges :: Edge e v => [e] -> Path e
pathFromEdges (x : xs) = foldl extend (pathFromEdge x) xs
pathFromEdges [] = error "pathFromEdges: Empty list of edges"
```

```haskell
ghci> pathFromEdges [("a", "b", 10), ("b", "c", 20)]
Path 30 [("b","c",20),("a","b",10)]
```

Here the (`String`, `String`, `Integer`) instance is used. Notice how the list is reversed in order to expose the target of the path at the head of the list.

The `Path` type can actually be thought of as an edge, since it has a source, a target, and a weight. The vertex type is the same as the underlying edge's vertex type.

```haskell
instance Edge e v => Edge (Path e) v where
  source (Path _ es) = source (last es)
  target (Path _ es) = target (head es)
  weight (Path w _)  = w
```

This instance is only possible if, just like edges, paths themselves can be compared for equality.

```haskell
ghci> weight (pathFromEdges [("a", "b", 10), ("b", "c", 20)])
30
```

Building on this abstraction of edges and vertices, a graph can be given by the following class interface:

There are many choices for describing a graph interface. This interface focuses on querying rather than constructing graphs.

```haskell
class Edge e v => Graph g e v | g -> e where
  vertices  :: g -> [v]
  edges     :: g -> [e]
  edgesFrom :: g -> v -> [e]
  edgesTo   :: g -> v -> [e]
  velem     :: v -> g -> Bool
  eelem     :: e -> g -> Bool
```

Usually it is best for the more stable parameters of a function to come first, so graphs before vertices. However, the "elem" functions usually have their arguments in the order so that they can be written as x `elem` xs, to mimic the mathematical notation $x \in X$.

Instances of this type class can be thought of as graphs: creating an instance for a graph takes three type parameters: the first indicates the representation of the graph, the second is for the representation of an edge, and the third is the representation of a vertex.

A simplistic representation of a graph is simply as a list of edges. This can certainly be achieved if the edges can be compared for equality.

```
instance (Eq e, Edge e v) => Graph [e] e v where
  vertices es = nub (map source es ++ map target es)
  edges es     = es
  edgesFrom es v = [ e | e <- es, v == source e ]
  edgesTo   es v = [ e | e <- es, v == target e ]
  velem v es = v `elem` vertices es
  eelem v es = v `elem` edges es
```

The nub function, from Data.List, removes duplicate elements

Thus, the following example is a legitimate graph:

```
example2 :: [(String, String, Integer)]
example2 = [("s","t",10), ("s","y",5), ("t","x",1),
            ("t","y",2),  ("y","t",3), ("y","x", 9),
            ("x","z",4),  ("z","x",6), ("y","z",2),
            ("z","s",7)]
```

It is possible to extract the vertices from this example:

```
ghci> vertices example2
["s","t","y","x","z"]
```


Figure 3: Graph of example2, bold lines indicate edges in the shortest path from *s* to other vertices.

The GameState type is an example of a graph: it holds information about planets, which are the vertices, and wormholes, which are the edges. The class instance shows how to extract this information.

```
instance Graph GameState (WormholeId, Wormhole) PlanetId where
  vertices (GameState ps _ _) = M.keys ps
  edges     (GameState _ ws _) = M.assocs ws
  edgesTo   st pId = M.toList (wormholesTo (Target pId) st)
  edgesFrom st pId = M.toList (wormholesFrom (Source pId) st)
  velem pId      (GameState ps _ _) = M.member pId ps
  eelem (wId, _) (GameState _ ws _) = M.member wId ws
```

This is not an efficient implementation, but it does allow for rapid prototyping, and querying the structure directly.

```
ghci> edgesFrom example1 0
[(0,Wormhole 0 1 1),(1,Wormhole 0 2 1),
 (2,Wormhole 0 3 1),(3,Wormhole 0 4 1)]
```

There are several representations of graphs that are possible and to this end, it is useful to work with a common interface for graphs, thus allowing experimentation with different implementations.

7   Shortest Paths

Players start with bases at some given distance away from each other. Knowing the distance to every planet helps to estimate how long it will be before there can possibly be any conflict there.

Working out the shortest distance through the galaxy can be reduced to an instance of Dijkstra's algorithm. The version explored here finds a list of the shortest paths from a root vertex to each other reachable vertex in the graph.

There are three main data structures at the heart of Dijkstra's algorithm: a graph, a priority queue, and a set. The graph is required to pull out edges for consideration when constructing the shortest path. The priority queue holds tentative shortest paths to vertices neighbouring the targets of shortest paths that have already been found. The set holds vertices that have yet to be explored.

Before working with Dijkstra's algorithm, it will be useful to understand priority queues.

7.1   Priority queues

A priority queue is a data structure that contains elements which have a given priority. Elements can be inserted into a priority queue, and elements can be extracted from the priority queue in ascending order.

The ordering imposed by a priority queue is given by a function which compares two elements and returns an `Ordering`. The values of type `Ordering` are one of `LT` (less than), `EQ` (equal), or `GT` (greater than). The `Util.hs` module has some functions that turn a binary operation that returns an ordering into a function that corresponds to the more traditional (<=) and (==) in Haskell.

```
lt :: (a -> a -> Ordering) -> (a -> a -> Bool)
lt cmp x y = cmp x y == LT

gt :: (a -> a -> Ordering) -> (a -> a -> Bool)
gt cmp x y = cmp x y == GT

lte :: (a -> a -> Ordering) -> (a -> a -> Bool)
lte cmp x y = cmp x y /= GT

eq :: (a -> a -> Ordering) -> (a -> a -> Bool)
eq cmp x y = cmp x y == EQ
```

A useful intuition is that a priority queue is abstractly simply an ordered list. Element order must be maintained when elements are added or removed from this list. The `PQueue` interface allows for different implementations of a priority queue:

Notice that some operations like `toPQueue`, `fromPQueue` and `detach` can be implemented in terms of other operations, so default implementations are provided for these. In case there is a more efficient implementation for a particular instance, then these defaults can be overridden (as you will see shortly).

```haskell
class PQueue pqueue where
  toPQueue   :: (a -> a -> Ordering) -> [a] -> pqueue a
  toPQueue cmp xs = foldr insert (empty cmp) xs

  fromPQueue :: pqueue a -> [a]
  fromPQueue = unfoldr detachMaybe

  priority :: pqueue a -> (a -> a -> Ordering)

  empty :: (a -> a -> Ordering) -> pqueue a
  isEmpty :: pqueue a -> Bool

  insert :: a -> pqueue a -> pqueue a

  extract :: pqueue a -> a
  discard :: pqueue a -> pqueue a
  detach  :: pqueue a -> (a, pqueue a)
  detach q = (extract q, discard q)

  detachMaybe :: pqueue a -> Maybe (a, pqueue a)
  detachMaybe q
    | isEmpty q = Nothing
    | otherwise = Just (detach q)
```

This intuition can be firmed up by giving laws that govern a valid priority queue instance:

These laws appeal to operations on lists, which can be found at `https://hackage.haskell.org/package/base-4.12.0.0/docs/Data-List.html`.

$$\text{fromPQueue (toPQueue cmp xs)} = \text{sortBy cmp xs} \tag{1}$$

$$\text{fromPQueue (empty cmp)} = \text{[]} \tag{2}$$

$$\text{isEmpty xs} = \text{nil (fromPQueue xs)} \tag{3}$$

$$\begin{aligned}\text{fromPQueue (insert x xs)} = \\ \text{insertBy (eq (priority xs)) (fromPQueue xs)}\end{aligned} \tag{4}$$

$$\text{extract xs} = \text{head (sortBy (priority xs) (fromPQueue xs))} \tag{5}$$

$$\text{discard xs} = \text{tail (sortBy (priority xs) (fromPQueue xs))} \tag{6}$$

$$\text{detach xs} = \text{(extract xs, discard xs)} \tag{7}$$

Notice that `toPQueue . fromPQueue` is not required to be the identity function: there is space for different priority queue representations of the same set of elements.

Perhaps the most natural implementation of a priority queue is a sorted list, together with the function that will determine the ordering of elements (i.e. the total ordering relation that the list is sorted by):

```haskell
data PList a = PList (a -> a -> Ordering) [a]
```

This type, and its `PQueue` instance, can be found in `PList.hs`. To witness that this is indeed a priority queue requires a `PList` instance of `PQueue`:

```
instance PQueue PList where
  toPQueue cmp xs = PList cmp (sortBy cmp xs)
  fromPQueue (PList _ xs) = xs
  empty cmp = PList cmp []
  isEmpty (PList _ xs) = null xs
  priority (PList cmp _) = cmp

  insert x (PList cmp []) = PList cmp [x]
  insert x ps@(PList cmp xs)
    | x <= y    = cons x ps
    | otherwise = cons y (insert x ys)
   where (<=) = lte cmp
         (y, ys) = detach ps
         cons x (PList cmp xs) = PList cmp (x:xs)

  extract (PList cmp (x:xs)) = x
  discard (PList cmp (x:xs)) = PList cmp xs
```

In the implementation of Dijkstra's shortest path algorithm, the priority queue holds paths, ordered by their total length. This ordering is implemented by the `cmpPath` function (given in `Graph.hs`).

```
cmpPath :: Path v -> Path v -> Ordering
cmpPath (Path d _) (Path d' _) = compare d d'
```

7.2   Dijkstra's Algorithm

The purpose of Dijkstra's algorithm is to find shortest paths. Rather than finding the shortest path from a source vertex to a particular target, the algorithm presented below finds shortest paths from a source vertex to all reachable vertices in the graph.

The `shortestPaths` function, provided in `ShortestPaths.hs`, operates on a given graph `g` and source vertex `v` to return a list of the shortest paths from `v`.

```
shortestPaths :: forall g e v. Graph g e v
              => g -> v -> [Path e]
shortestPaths g v = dijkstra g (vertices g \\ [v]) ps where
  ps :: PList (Path e)
  ps = toPQueue cmpPath (map pathFromEdge (edgesFrom g v))
```

The main workhorse of `shortestPaths` is the `dijkstra` function, which implements Dijkstra's algorithm. The algorithm operates on a graph (`g` here), and maintains a list of nodes yet to be visited. Here, this list is

initially all the vertices of `g` except for the starting vertex `v`. Finally, the algorithm also maintains a list of shortest path candidates as a priority queue of paths (`ps`), which here is initially the edges starting at `v`.

Notice that the `dijkstra` function does not take in the initial vertex as an argument. Instead, it operates on the priority queue of candidate shortest paths and grows them in a minimal way. The `shortestPaths` function takes care of calling `dijkstra` with all the edges from the starting node.

Here's how the algorithm works:

1. If the list of unvisited nodes or the priority queue of candidate edges is empty, then the algorithm is finished and returns the empty list.

2. Otherwise, the algorithm selects the minimum path `p` from the priority queue `ps` (remember that this can simply be achieved by using `detach`.). In the simple case, the target vertex `t` of the path `p` has already been visited, so a shortest path has already been found to that node. The algorithm therefore proceeds with the remaining priority queue that does not include `p`.

3. If the target vertex `t` of the path `p` is in the list of unvisited vertices in `us`, then the path `p` is added to the list of solutions: this is the shortest path to `t`. The remaining solutions are then found by recursively calling the algorithm with an updated set of unvisited nodes and an updated priority queue. The set of unvisited nodes is updated to remove `t`. Since `p` is a shortest path, new candidate shortest paths are those that extend `p` by the edges from `t`.

The algorithm thus iteratively builds a list of all shortest paths starting from the source. The correctness of this algorithm (that it indeed returns the shortest paths) can be proved by induction on the length of the visited nodes. The second step above relies on the invariant that once a node has been visited, we know the shortest path to it. The third step maintains this invariant. The key idea is that the shortest candidate path is always guaranteed to be shorter than all the other paths, even the ones not yet included in the candidates. This is because the list of candidates is grown by extending the existing candidates. Crucially, we assume that there are no negative edges, which means that any new candidate must be longer than the current shortest candidate.

Once again, the `dijkstra` function is highly polymorphic, and it works against any graph and priority queue implementation

---

Problem 6: *Dijkstra's algorithm*

Finish the implementation of `dijkstra` in `ShortestPaths.hs` by completing the definitions for `us'` and `ps''`. `us'` is the new list of unvisited nodes to consider in the recursive call, and `ps''` is the updated

priority queue `ps'` with the new candidates inserted into it. This case corresponds to step 3 from the above description.

```
dijkstra :: forall g e v pqueue.
  (Graph g e v, PQueue pqueue) =>
  g -> [v] -> pqueue (Path e) -> [Path e]
dijkstra g [] ps = []
dijkstra g us ps
  | isEmpty ps  = []
  | t `elem` us =
      let us' :: [v]
          us' = undefined
          ps'' :: pqueue (Path e)
          ps'' = undefined
      in p : dijkstra g us' ps''
  | otherwise  = dijkstra g us ps'
  where
    (p, ps') = detach ps
    t = target p
```

HINT: You might find `foldr` useful for updating the shortest path candidates.

The shortest path between any two planets in `example1` can be computed by simply calling `shortestPaths`, since `GameState` is a graph.

```
ghci> shortestPaths example1 0
[ Path 1 [(0,Wormhole 0 1 1)]
, Path 1 [(1,Wormhole 0 2 1)]
, Path 1 [(2,Wormhole 0 3 1)]
, Path 1 [(3,Wormhole 0 4 1)] ]
```

This is a rather uninformative example all of the planets are connected to the source and not to each other.

A more useful try is with `example2`, and this also works with `shortestPaths` because lists of edges are graphs:

```
ghci> shortestPaths example2 "s"
[ Path 5 [("s","y",5)]
, Path 7 [("y","z",2),("s","y",5)]
, Path 8 [("y","t",3),("s","y",5)]
, Path 9 [("t","x",1),("y","t",3),("s","y",5)] ]
```

Although the implementation here is correct (assuming your solution to `dijkstra` is correct!), one problem with the code is that it is very inefficient. This is largely due to the incorrect choice of data structures:

the graph representation, the priority queue of paths, and the set of unvisited vertices all use basic structures that have suboptimal complexities. For instance, while extracting from the priority queue is a constant time operation, insertion takes linear time in the length of the queue, because in order to maintain the invariant that the elements are ordered, the insertion function walks through the elements to find the correct location for the new element.

7.3   A Heap of Paths

The first optimisation will be to replace the `PList` data structure with a more efficient one. There is a myriad of data structures to choose from, and picking the right one can be challenging. It is always a good idea to consider how the data structure is being used, and which operations need to be efficient, as there are tradeoffs where an operation might be more efficient at the expense of another one. For our priority queue, we need efficient extraction of the minimum element and efficient insertion, as these are the two operations used in Dijkstra's algorithm.

To satisfy these requirements, we will use a *binary heap* data structure. At its simplest, this is a binary tree that maintains the minimum element at the root of the tree, supporting constant-time access to the smallest element. The specific flavour we will be building is a *leftist heap*, which supports efficient merging. The data structure for a heap, found in `Heap.hs`, is as follows:

```haskell
data Heap a = Heap (a -> a -> Ordering) (Tree a)
data Tree a = Nil | Node Int (Tree a) a (Tree a)
```

The `Heap` type packages a comparison function with a binary tree.

The `Tree` datatype is a binary tree with values at the nodes, but in addition, each node also contains an `Int`, called the *rank* of the tree.

```haskell
rankTree :: Tree a -> Int
rankTree Nil           = 0
rankTree (Node h l x r) = h

rankHeap :: Heap a -> Int
rankHeap (Heap _ t) = rankTree t
```

The rank of a tree is the distance from the root node to the nearest leaf node. This can of course always be computed on the fly, but for efficiency, the `Node` constructor caches this value. Furthermore, the rank of the left subtree is always at least as high as that of the right subtree (which is where the name *leftist* comes from).   The following *smart constructor* ensures that when building a new root, the rank and the above invariant are maintained:

A *smart constructor* is a function that acts as a drop-in replacement for a constructor, while maintaining the invariants of the data structure. Here, using `node` instead of `Node` will make sure that the rank invariant always holds.

```haskell
node :: Tree a -> a -> Tree a -> Tree a
node l x r
  | hl < hr   = Node (hl + 1) r x l
  | otherwise = Node (hr + 1) l x r
 where
  hl = rankTree l
  hr = rankTree r
```

The heap operations will ensure that the smallest element is always at the root of the tree. In other words, for any `Node h l x r`, x is smaller than l and r (and recursively this means that x is smaller than every element in the tree).

The first operation on heaps is `mergeHeap` which merges two heaps together such that the resulting heap will contain the elements of both heaps while maintaining the invariants. `mergeHeap` simply calls `mergeTreee` with the comparison functions of the first heap (which should be the same as that of the second heap).

Only heaps that have been constructed with the same comparison function should be merged. It is not possible in Haskell to check functions for equality, so this invariant can not be checked, but it is assumed throughout.

```haskell
mergeHeap :: Heap a -> Heap a -> Heap a
mergeHeap (Heap cmp l) (Heap _ r) = Heap cmp (mergeTree cmp l r)
```

Problem 7: *Merging trees*

Implement the `mergeTree` operation which merges two leftist trees.

```haskell
mergeTree  :: (a -> a -> Ordering)
           -> Tree a -> Tree a -> Tree a
mergeTree cmp l r = undefined
```

If either tree is empty, the result should be the other tree. When merging two nodes, first find out which tree should be merged into which (such that the root is always the smallest element in the tree), then recursively merge one tree into the right subtree of the other, potentially swapping at the end if the resulting right tree is taller than the left. Merging into the right (shorter) subtree ensures that the tree never gets too unbalanced. This in turn means that we will always have a logarithmic upper bound on the depth at each merge. HINT: use the `node` smart constructor to maintain the rank invariant.

Problem 8: *Heap Operations*

Provide the instance of `PQueue` that shows how a heap can be used as a priority queue, by filling out the following definitions:

```haskell
instance PQueue Heap where
  priority :: Heap a -> (a -> a -> Ordering)
```

```
   priority = undefined

   empty :: (a -> a -> Ordering) -> Heap a
   empty p = undefined

   isEmpty :: Heap a -> Bool
   isEmpty = undefined

   insert :: a -> Heap a -> Heap a
   insert = undefined

   extract :: Heap a -> a
   extract = undefined

   discard :: Heap a -> Heap a
   discard = undefined

 HINT: insertion can be thought of as a special case of merging.
```

The definition of `shortestPaths'`, given in `ShortestPaths.hs`, uses a `Heap` instead of a `PList` for the underlying priority queue representation.

```
shortestPaths' :: forall g e v . Graph g e v
               => g -> v -> [Path e]
shortestPaths' g v = dijkstra g (vertices g \\ [v]) ps where
  ps :: Heap (Path e)
  ps = foldr (insert . pathFromEdge) (empty cmpPath) (edgesFrom g v)
```

This is exactly the same code as for `shortestPaths`, except that the type annotation for `ps` specifies that it should use a `Heap`.

7.4   Adjacency List Graph

The graph representation used so far has been the simple `GameState` instance. The problem here is that it is not very efficient for the `dijkstra` algorithm. There, the `edgesFrom` function is used to find all the edges from a vertex. This operation happens each time a shortest path has been found. In the instance of `Graph` for `GameState`, the implementation filters through the list of all the wormholes.

A different representation of a graph is known as the *adjacency list*. An adjacency list is a list of pairs containing each vertex paired with all edges from that vertex. The type can be given by `AdjList`, provided in `AdjList.hs`:

```
newtype AdjList e v = AdjList [(v, [e])]
```

In other words, each pair (v, es) in the adjacency list is a vertex v, and each edge e in the edges es is such that source e = v.

Problem 9: *Adjacency List Graphs*

Implement an adjacency list representation that supports efficient lookup of the edges from a planet by filling in the following interface.

```
instance (Eq e, Edge e v) =>
    Graph (AdjList e v) e v where
  vertices  (AdjList ves)   = undefined
  edges     (AdjList ves)   = undefined
  edgesFrom (AdjList ves) s = undefined
  edgesTo   (AdjList ves) t = undefined
  velem v   (AdjList ves)   = undefined
  eelem e   (AdjList ves)   = undefined
```

HINT: You may use the standard list operations given in `Data.List`. You may also find list comprehensions useful.

7.5   Conflict Zones

Having established an algorithm for finding the shortest paths from a root vertex, it is possible to calculate the planets which each player can reach first. This can be achieved by running the shortest path algorithm from each player home base. As a final task, you should write an algorithm that calculates this for a given GameState.

Problem 10: *Conflict Zones*

In ShortestPaths.hs, define conflictZones, where conflictZones st p q takes in the game state st, and two planet IDs p and q, and returns a triple (ps, pqs, qs) where ps are the identities of planets that can be reached by p first, qs are the identities of planets that can be reached by q first, and pqs are the identities of planets that can be reached by both at the same time.

```
conflictZones :: GameState -> PlanetId -> PlanetId
              -> ([PlanetId], [PlanetId], [PlanetId])
```

Note that this should only be a function of the topology of the game: it does not need to take into account the number of ships garrisoned on planets, whether conquering is needed, the positions of fleets that are in flight or any other details.